

# Small Data

CMPT 732, Fall 2017

Lots of data sets aren't big. In fact, most aren't.

Modern phones have 4 GB of memory: if you have less than that, it must be “small”. Why use Spark for everything?

Even running locally, Spark has a  $\approx 10$  s startup time: any work that takes less than that makes **no** sense in Spark.

Good reasons to use Spark: [editorial content]

- You actually have big data.
- You think your data might be big in the future, and need to be ready.
- You have “medium” data and the startup time pays off when running locally on multiple cores.

## Spark for ETL

A very good use-case for Spark: ETL work that makes big data small.

Use Spark to extract/aggregate the data you really want to work with. Realize that made your data “small”. Move to some small data tools...

## Python Data Tools

Python is one of the most common choices for data science work. (The other is R.)

As a result, there are many very mature data manipulation tools in Python. You should know they exist.

## NumPy

Python's built-in data structures are not very memory-efficient: Python object overhead, references cause bad memory locality, etc.

Data you have will often have fixed types and sizes: exactly what C-style arrays are good at. [NumPy](#) provides efficient, typed arrays for Python.

```
import numpy as np
a = np.zeros((10000, 5), dtype=np.float32)
print(a)
print((a + 6).sum())
```

```
[[ 0.  0.  0.  0.  0.]
 [ 0.  0.  0.  0.  0.]
 [ 0.  0.  0.  0.  0.]
 ...
 [ 0.  0.  0.  0.  0.]
 [ 0.  0.  0.  0.  0.]
 [ 0.  0.  0.  0.  0.]]
300000.0
```

NumPy can do lots of manipulation on arrays (at C-implemention speeds). e.g.

- basic arithmetic
- datetime manipulation
- matrix/linear algebra operations
- sorting, searching

## Pandas

[Pandas](#) provides a DataFrame class for in-memory data manipulation. Pandas DataFrame  $\neq$  Spark DataFrame, but concepts are similar.

```
import pandas as pd
cities = pd.read_csv('cities.csv')
print(cities)
```

	city	population	area
0	Vancouver	2463431	2878.52
1	Calgary	1392609	5110.21
2	Toronto	5928040	5905.71
3	Montreal	4098927	4604.26
4	Halifax	403390	5496.31

Similar operate-on-whole-DataFrame API. Slightly different operations. Not lazily evaluated.

```
cities['area_m2'] = cities['area'] * 1e6
print(cities)
```

	city	population	area	area_m2
0	Vancouver	2463431	2878.52	2.878520e+09
1	Calgary	1392609	5110.21	5.110210e+09
2	Toronto	5928040	5905.71	5.905710e+09
3	Montreal	4098927	4604.26	4.604260e+09
4	Halifax	403390	5496.31	5.496310e+09

Pandas Series (==columns) are stored as NumPy arrays, so you can use NumPy functions if you need to.

```
print(type(cities['population'].values))
print(cities['population'].values.dtype)
```

```
<class 'numpy.ndarray'>
int64
```

A Pandas DataFrame can be converted to a Spark DataFrame:

```
cities_pd = pd.read_csv('cities.csv')
cities_spark = spark.createDataFrame(cities_pd)
cities_spark.show()
```

```
+-----+-----+-----+
|   city|population|   area|
+-----+-----+-----+
|Vancouver|  2463431|2878.52|
|  Calgary|  1392609|5110.21|
|  Toronto|  5928040|5905.71|
|Montreal|  4098927|4604.26|
|  Halifax|   403390|5496.31|
+-----+-----+-----+
```

... and a Spark DataFrame to Pandas **if** it will fit in memory in the driver:

```
cities_pd2 = cities_spark.toPandas()
print(cities_pd2)
```

	city	population	area
0	Vancouver	2463431	2878.52
1	Calgary	1392609	5110.21
2	Toronto	5928040	5905.71
3	Montreal	4098927	4604.26
4	Halifax	403390	5496.31

This is faster in Spark  $\geq 2.3$  if you use [the Apache Arrow option](#).

With NumPy and Pandas, you can do a lot of basic data manipulation operations.

They will likely be faster on small (and medium?) data: no overhead of managing executors or distributing data, but single-threaded.

## SciPy

The [SciPy](#) libraries include many useful tools to analyze data. Some examples:

- NumPy and Pandas
- Fourier Transforms (`scipy.fftpack`)
- Signal Processing (`scipy.signal`)
- Linear Algebra (`scipy.linalg`)
- Statistics (`scipy.stats`)
- Image processing (`scipy.ndimage`)
- Plots (`matplotlib`)

If those aren't enough, there are [SciKits](#) containing much more. e.g.

- Image processing (`scikit-image`)
- Video processing (`scikit-video`)
- Bioinformatics (`scikit-bio`)

## SciKit-Learn

[Scikit-learn](#) is probably going to be useful to you some time: implementations of many machine learning algorithms for Python (and NumPy-shaped data).

Compared to `pyspark.ml`: older and more battle-tested; [includes algorithms](#) that don't distribute well; doesn't do distributed computation.

## Python Libraries

Maybe the biggest pro-Python argument: it's used for data science and many other things, so libraries you need are implemented in Python.

[PyPI](#) is the package repository for Python. You can install packages with the `pip` command.

```
pip3 install --user scikit-learn scipy pandas
```

